

# Reinforcement Learning Project: Continuous Cartpole

---

**Done by:**

Simon Ging, Abdelrahman Younes, Naya Baslan

**Supervised by:**


Neurorobotics Lab, Uni Freiburg

# Agenda

- **Overview**
  - Choice of Algorithm
  - Performance analysis
- **Details**
  - Performance visualization
  - Project process
  - Algorithm details and improvements
  - Reward function analysis
  - Success metrics discussion
  - Visual comparison of result differences
  - Epsilon scheduler analysis
  - Runtime Analysis
  - HPO configuration space
  - All results in one plot
  - Best agent showcase
  - Qualitative evaluation of 2 other environments
- **Outlook**
- **References**

# Choice of Implemented Algorithm

- Choice: DQN<sup>[1]</sup> (discrete action space)
  - Popular algorithm, can solve Atari games
  - Lots of possible improvements<sup>[3][4][5]</sup>
  - Agents with discretized actions can solve given problem



$$\text{DQN Target: } Y_t^{\text{DoubleQ}} \equiv R_{t+1} + \gamma Q(S_{t+1}, \arg\max_a Q(S_{t+1}, a; \theta_t); \theta'_t)$$

Figure 1: Schematic diagram of DQN Algorithm<sup>[1]</sup> with Double Q Learning<sup>[3]</sup> and Dueling Networks<sup>[4]</sup>

# Performance Analysis

Experiment	K steps	Episodes	Evaluation (1000 ep)	Train score (last 100 ep)	Success metric	Converged / Total
Baseline	688.07±161.59	2851.13±714.40	303.42±177.98	191.71±126.34	Eval ≥ 400 (every 20K steps)	<b>10/15</b>
Best (eval @5K)	<b>189.02±28.41</b>	<b>879.73±80.56</b>	414.96±18.02	59.81±45.44	Eval ≥ 400 (every 5K steps)	<b>15/15</b>
Best (eval @20K)	197.41±35.29	893.93±110.19	414.85±12.16	81.16±57.91	Eval ≥ 400 (every 20K steps)	<b>15/15</b>
Best (train score)	512.08±162.60	1521.27±210.33	<b>420.25±3.10</b>	<b>400.42±0.75</b>	Score of last 100 episodes ≥ 400	<b>14/15</b>
Best (shaped reward)	256.03±54.99	1027.00±163.02	409.688±40.87	13.23K±71.80K	Eval ≥ 400 (every 5K steps)	<b>15/15</b>

Table 1: Quantitative Analysis of 5 selected experiments

- Careful HP optimization improved the results
- Different success metric have their advantages
- High variance, repeated experiments needed
- Highly time efficient
  - ~5min GPU runtime for experiment *Best (eval @5K)*

# Performance visualization

Baseline improved significantly, eval success metric outperforms train metric

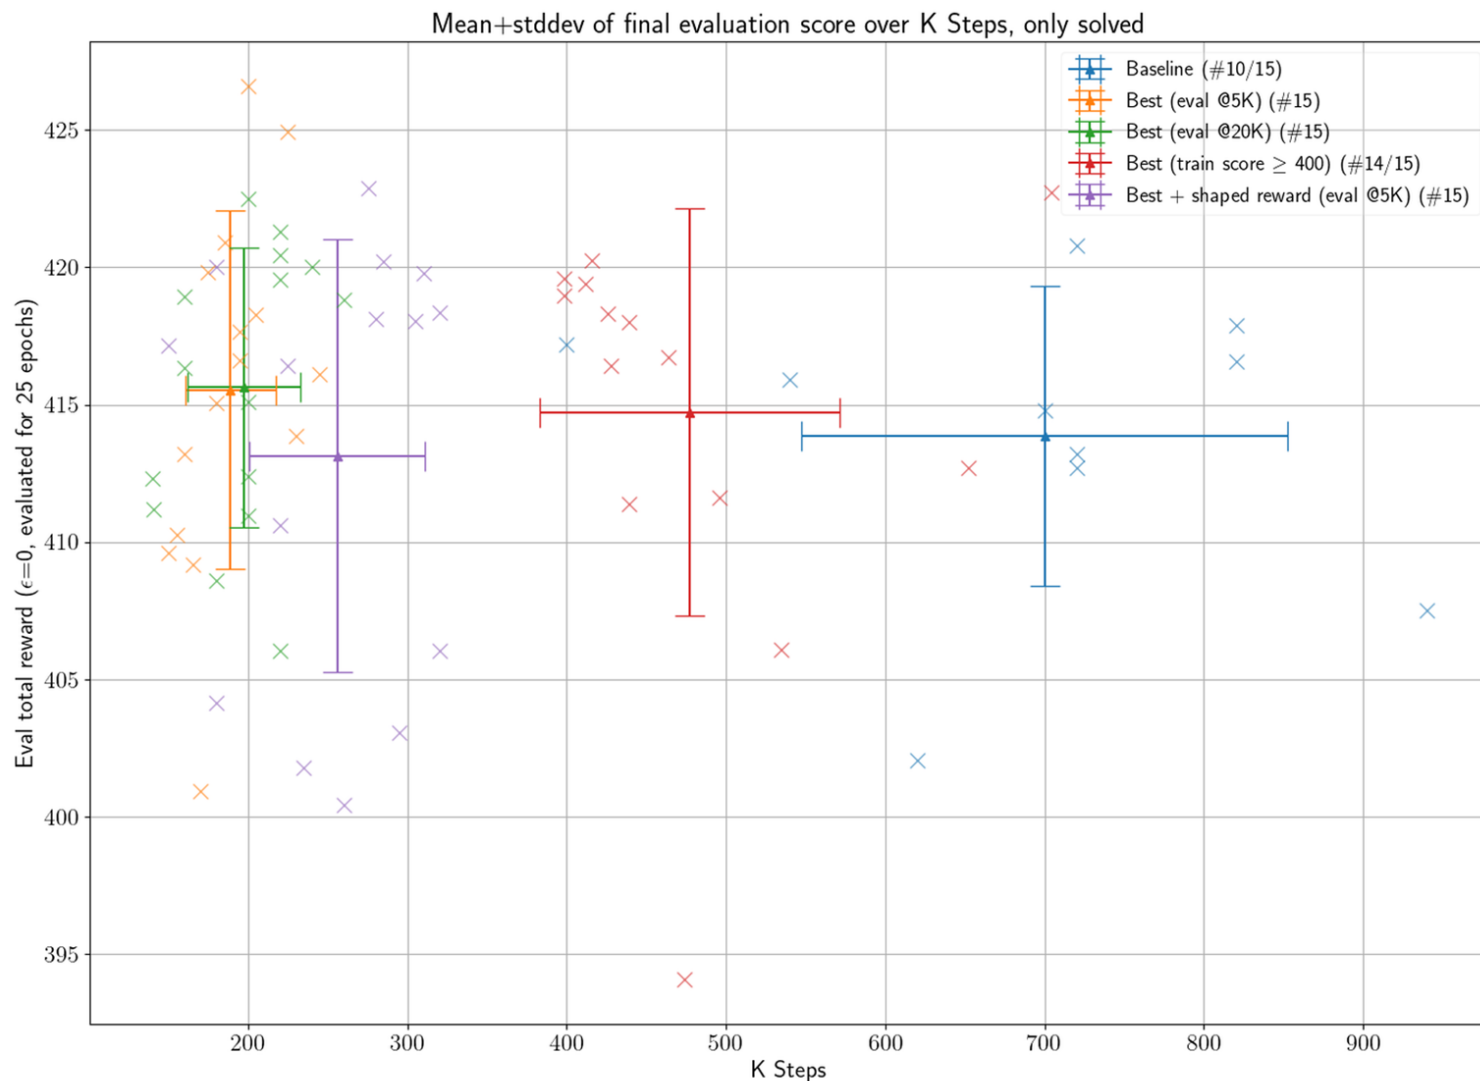


Figure 2: Convergence speed comparison of 5 selected experiments

# Project Process Details

- Baseline: DQN with local and target network<sup>[1]</sup>
  - Using soft target network update<sup>[2]</sup>
  - Converges after ~700k steps, but unstable
- Experiment setup
  - High variance: At least 5 runs per experiment
  - Default success metric: Eval score  $\geq 400$ 
    - $\epsilon=0$ , 25 episodes, evaluated every 20K steps
- Algorithm improvements
  - Double Q Learning<sup>[3]</sup>
  - Dueling Network Architecture<sup>[4]</sup>
- Local search for better hyperparameters
- In-depth evaluation

- Experiences are stored in a buffer and randomly sampled
- Uses two Networks
  - Online net predicts state-action-values  $y^{\circ} = Q(s_t, a_t | \phi^{\circ})$
  - Target net predicts targets for TD Error with weights  $\phi^{\top}$
- Every  $K=4$  steps:
  - Update  $\phi^{\circ}$  with gradient of  $\text{MSELoss}(\text{prediction}, \text{target})$
  - Soft update<sup>[2]</sup>  $\phi^{\top} = (1 - \tau) \phi^{\top} + \tau \phi^{\circ}$
- Act  $\epsilon$ -greedily w.r.t  $Q(s_t, a_t | \phi^{\circ})$ 
  - Decay  $\epsilon$  during training, evaluate with low or zero  $\epsilon$

# DQN Improvements - Double Q Learning

- DQN loss<sup>[1]</sup> 
$$L_i(\theta_i) = \mathbb{E}_{s,a,s',r \sim D} \left( \underbrace{r + \gamma \max_{a'} Q(s', a'; \theta_i^-)}_{\text{target}} - Q(s, a; \theta_i) \right)^2$$
- DQN target<sup>[3]</sup> 
$$Y_t^{\text{DQN}} \equiv R_{t+1} + \gamma \max_a Q(S_{t+1}, a; \theta_t^-)$$
  - Target network both chooses and evaluates target action
- Double Q-learning<sup>[3]</sup> target
$$Y_t^{\text{DoubleQ}} \equiv R_{t+1} + \gamma Q(S_{t+1}, \underset{a}{\operatorname{argmax}} Q(S_{t+1}, a; \theta_t); \theta'_t)$$
  - Choose best action for target with online network
  - Estimate its value with target network



# DQN Improvements - Dueling Networks

- DQN<sup>[1]</sup>: predict  $Q(s, a)$  directly
- Dueling Networks Architecture<sup>[4]</sup>:
  - Predict state-value  $V$  and advantage function  $A$ :

- Compute Q values:

$$Q(s, a; \theta, \alpha, \beta) = V(s; \theta, \beta) + \left( A(s, a; \theta, \alpha) - \frac{1}{|\mathcal{A}|} \sum_{a'} A(s, a'; \theta, \alpha) \right)$$

- Network graphs with image input
- We replaced convs with fully-connected layers
- We use a single hidden layer instead one for each  $V$  and  $A$




Figure 3: DQN vs Dueling Networks<sup>[4]</sup>

# Reward Function Analysis




Figure 4: Visualization of sparse and shaped reward

- Environment has sparse rewards
- Shaped reward tested as above
  - Agent performance dropped significantly
- Final configuration uses sparse reward

# Different success metrics

- High evaluation result was our goal
  - Evaluate every 20k (or 5k) steps for 25 episodes
  - Mean score  $\geq 400$ : Success
  - Sparse reward - independent of shaped reward choice
  - $\epsilon = 0$  - independent of exploration strategy
  - Overfits the agent to the problem
- Train metric for comparison
  - Mean train reward of last 100 episodes  $\geq 400$ : Success
  - Must be tuned for shaped reward
  - Harder to achieve / can be unstable
  - Probably needs more fine-tuning of  $\epsilon$ -decay
  - Very dependent on optimization process
  - Agents obtained this way are performing slightly better

# Best config, eval metric, evaluation score

Algorithm is stable, all 15 runs converged after 250K steps




Figure 5: Evaluation curve of best config with eval metric

# Best config, train metric, evaluation score

Train score metric is unstable. Too high  $\epsilon$ ?

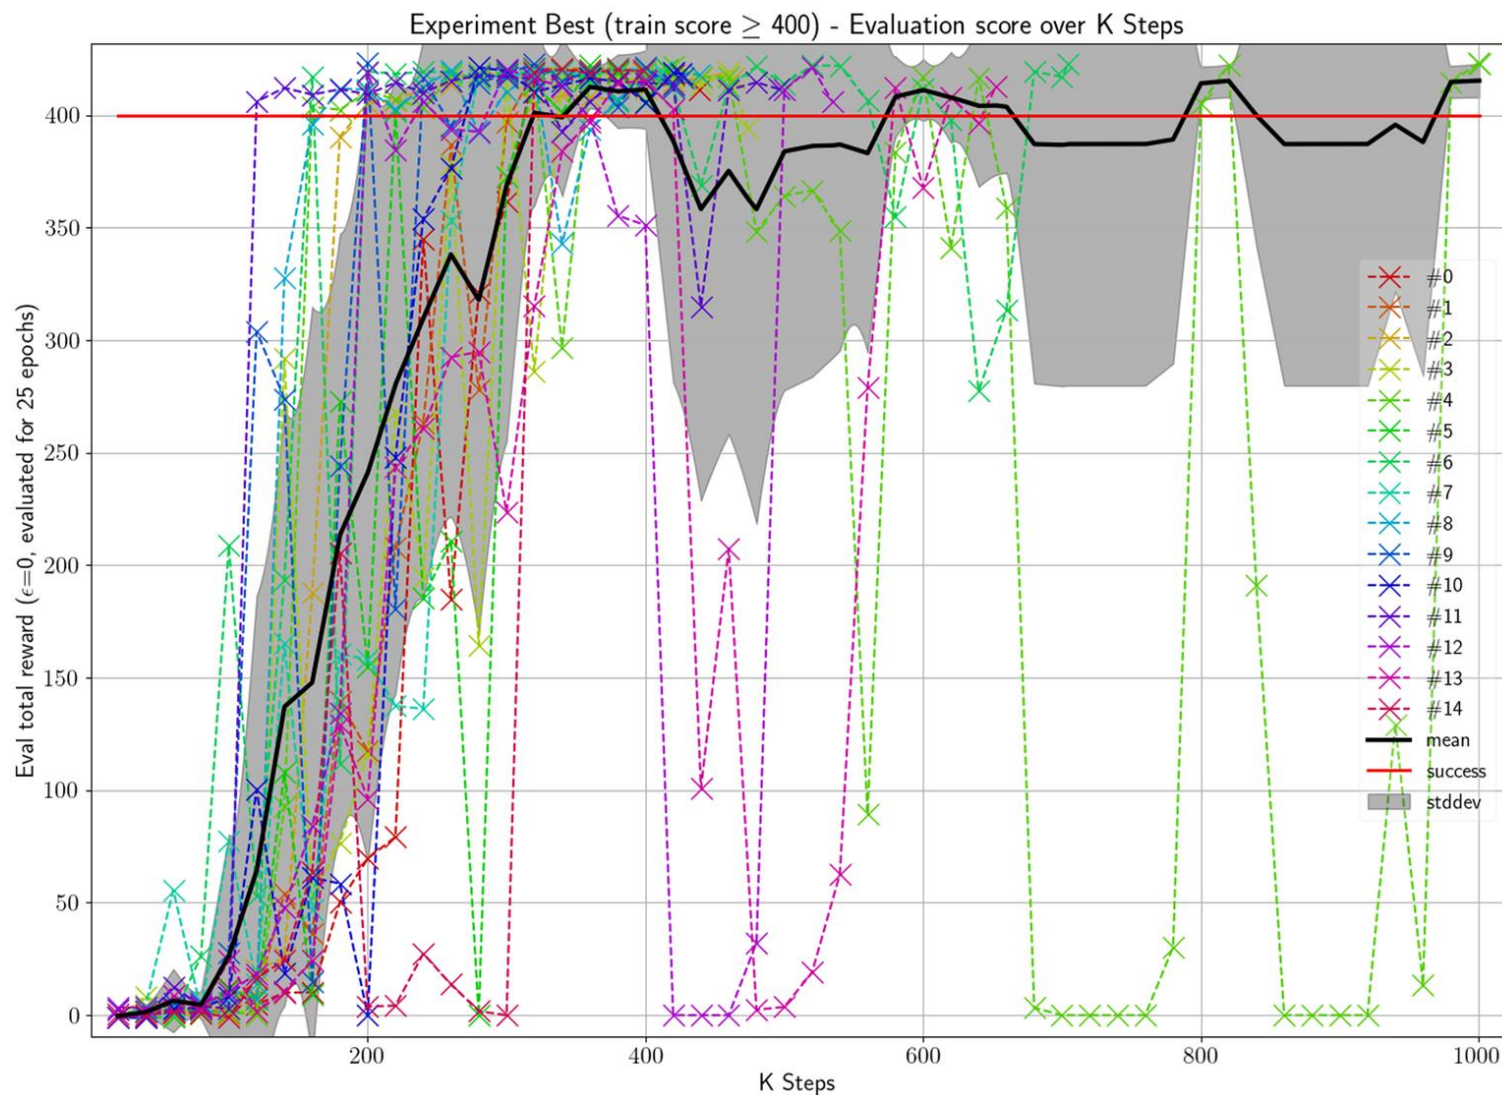


Figure 6: Evaluation curve of best config with train metric



# Best config, eval metric, training score

Agents converge before training score is high

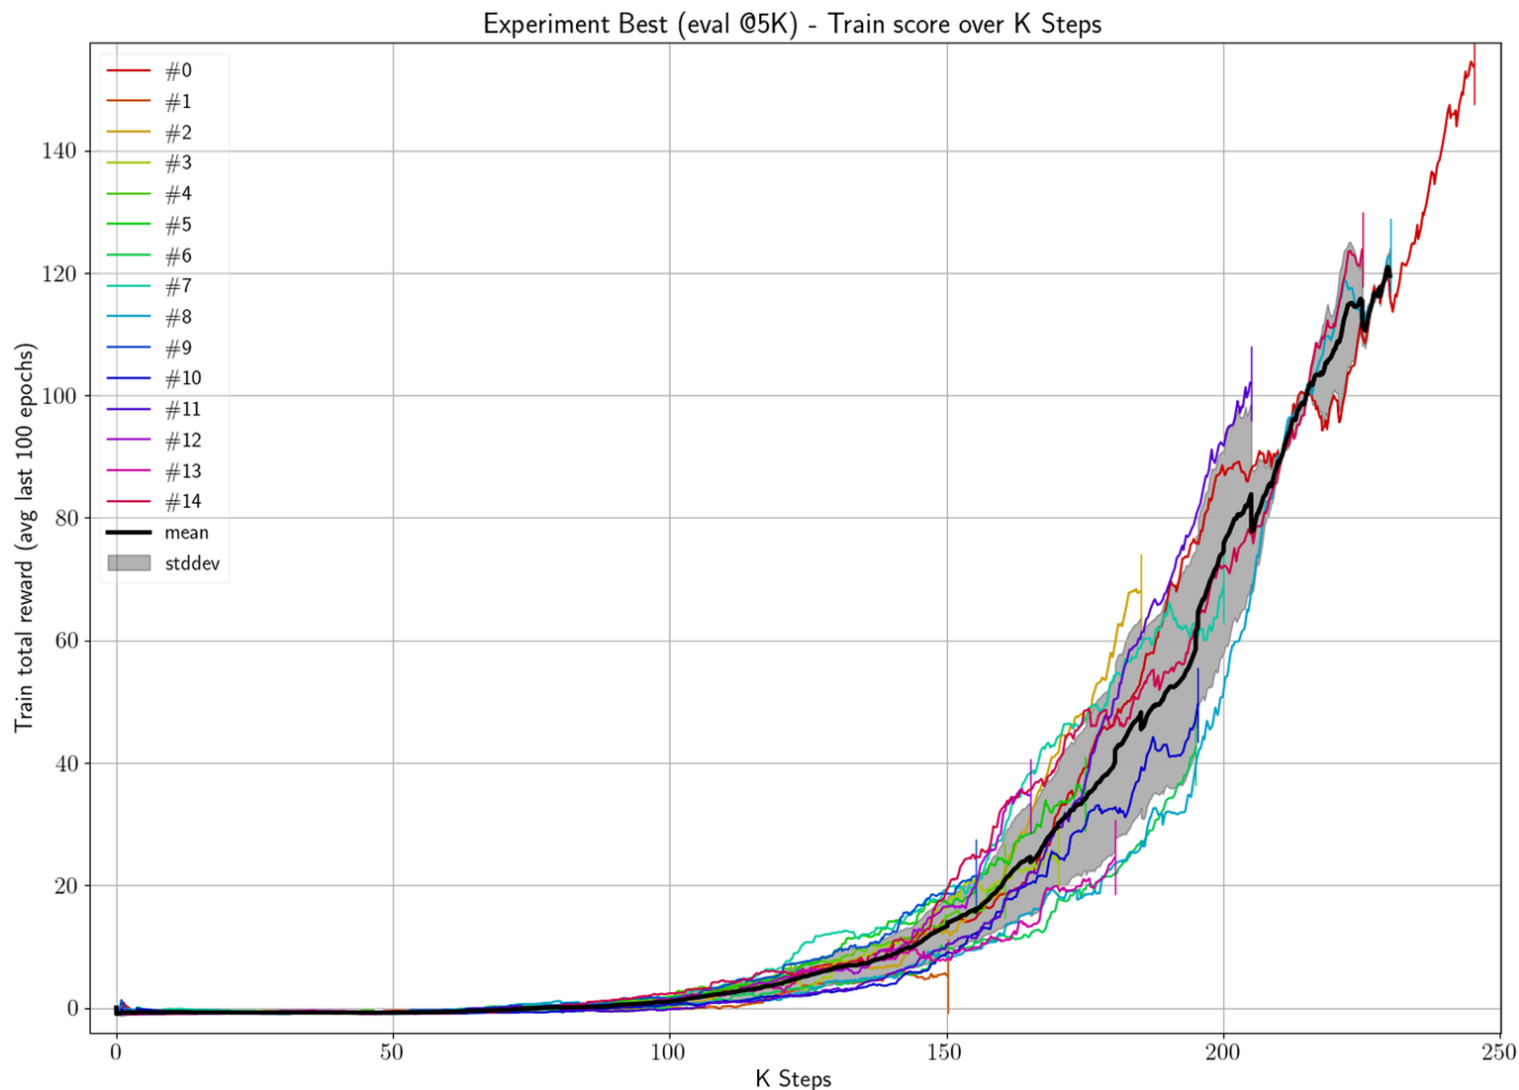


Figure 7: Training curve of best config with eval metric

# Best config, train metric, training score

Reaching  $\geq 400$  train reward over last 100 episodes seems difficult

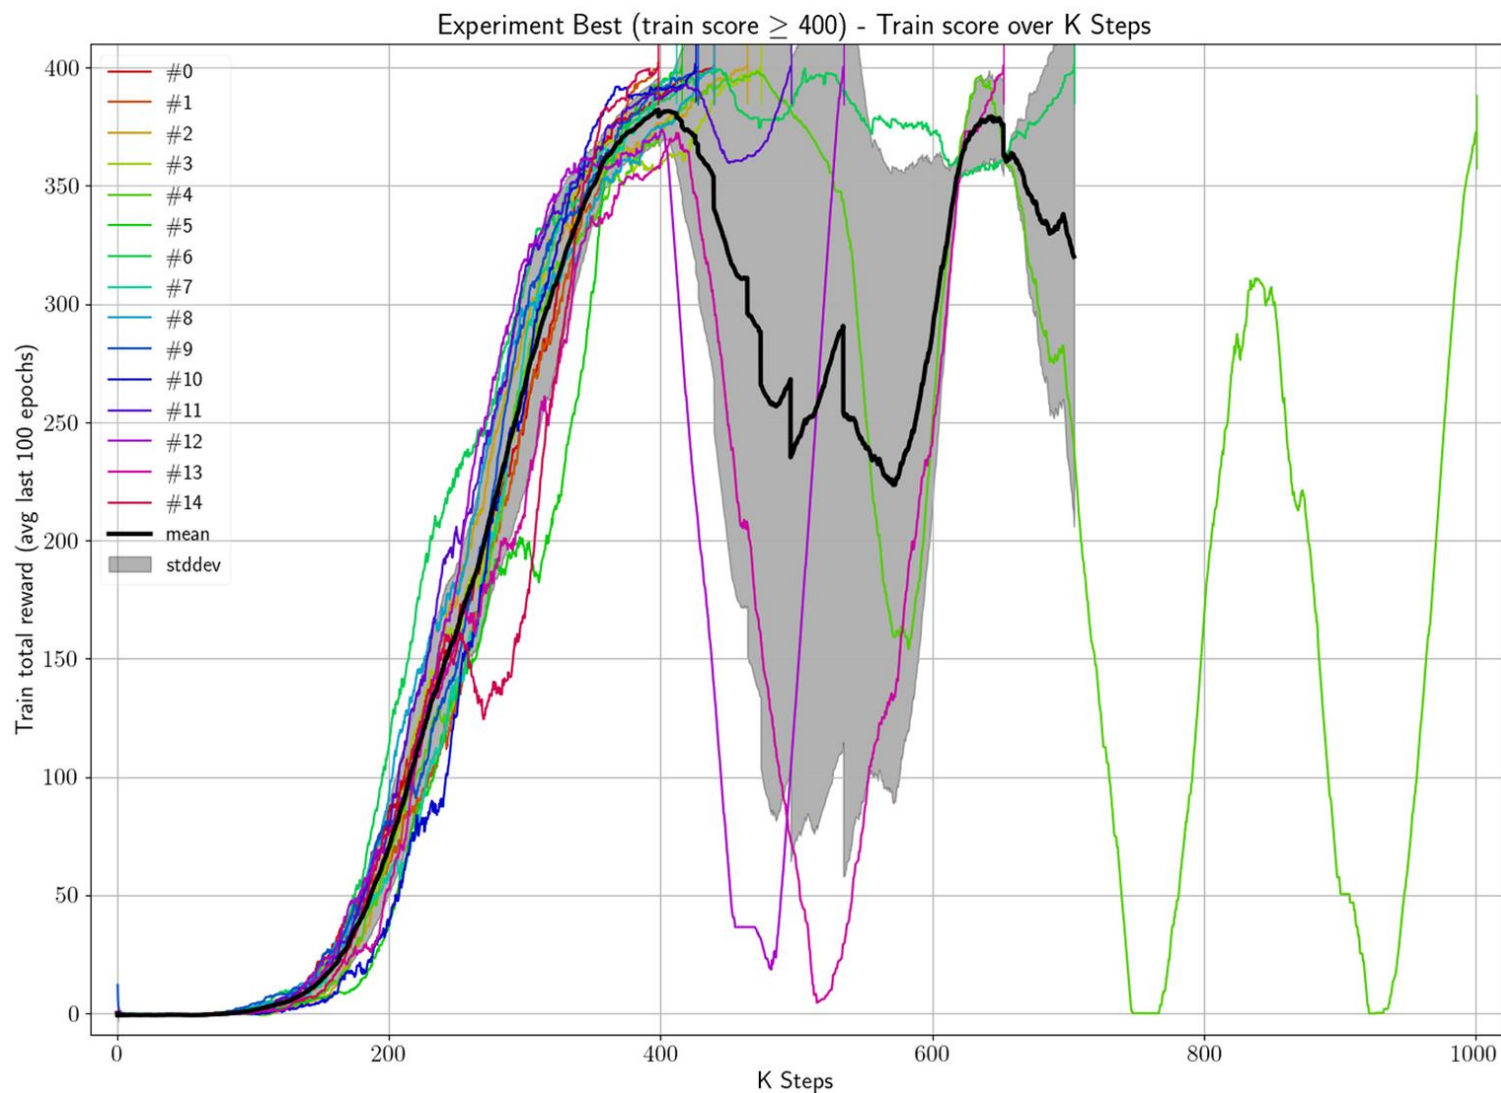


Figure 8: Training curve of best config with train metric

# Best config, eval metric, quality test

Agents are very good but sometimes still fail (tested for 1000 episodes)

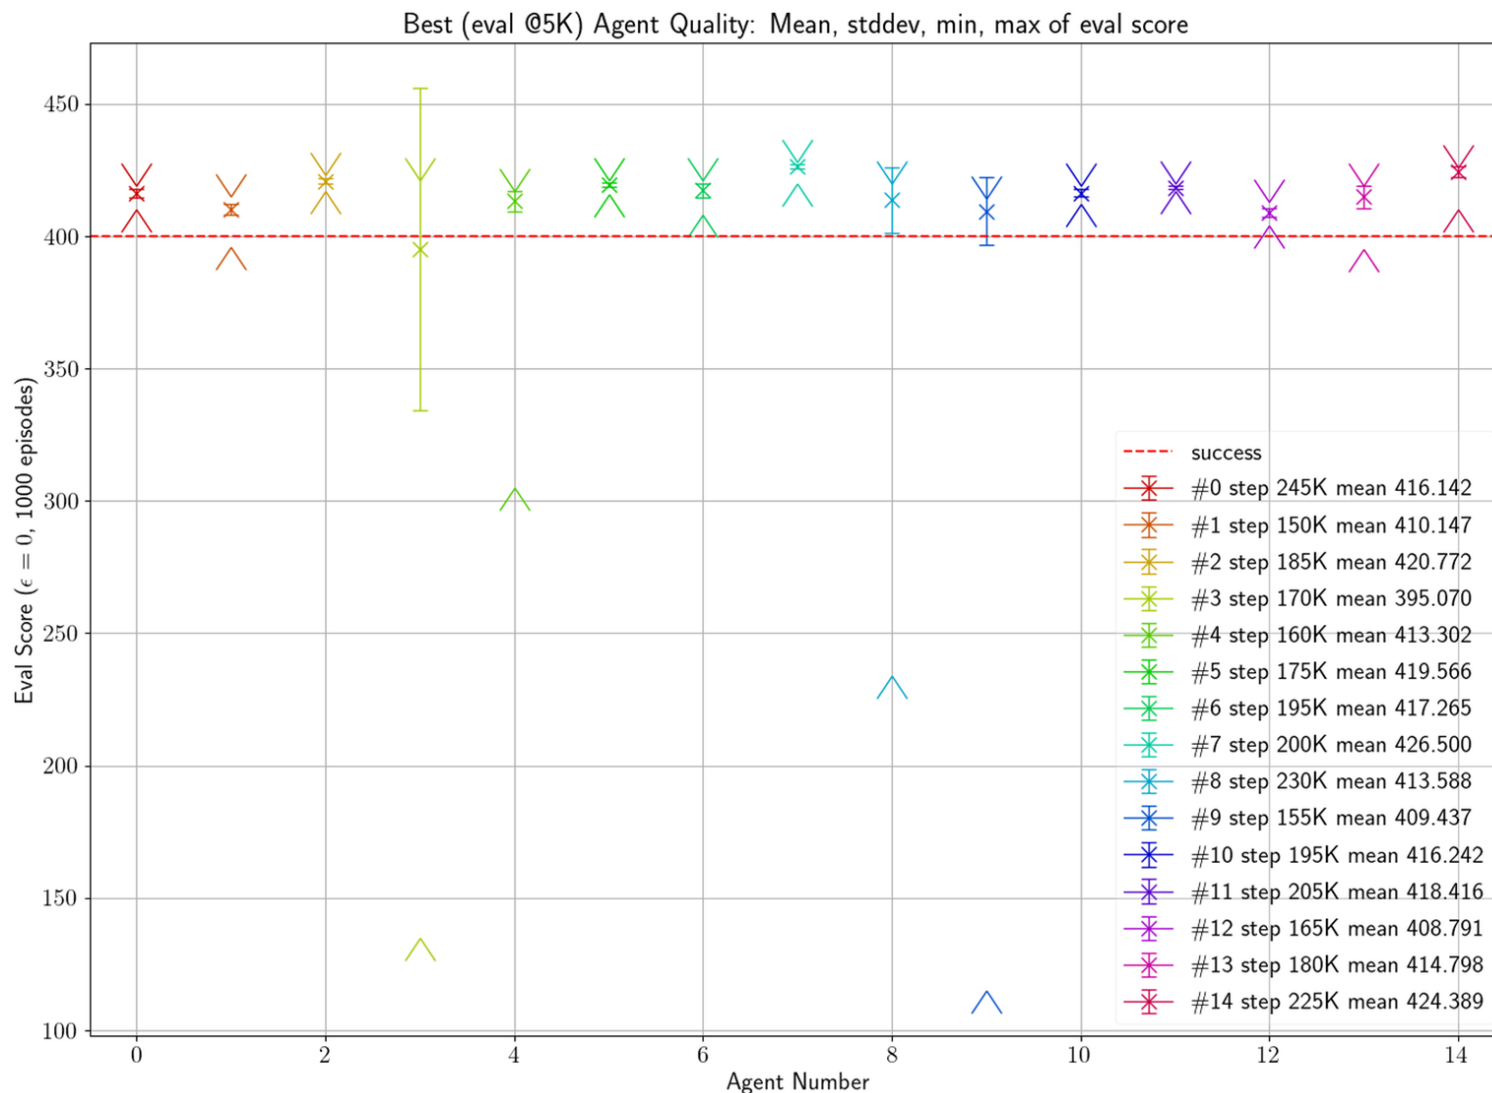


Figure 9: Quality test of best config with eval metric



# Best config, train metric, quality test

Train metric is unstable and slower but provides much better agents

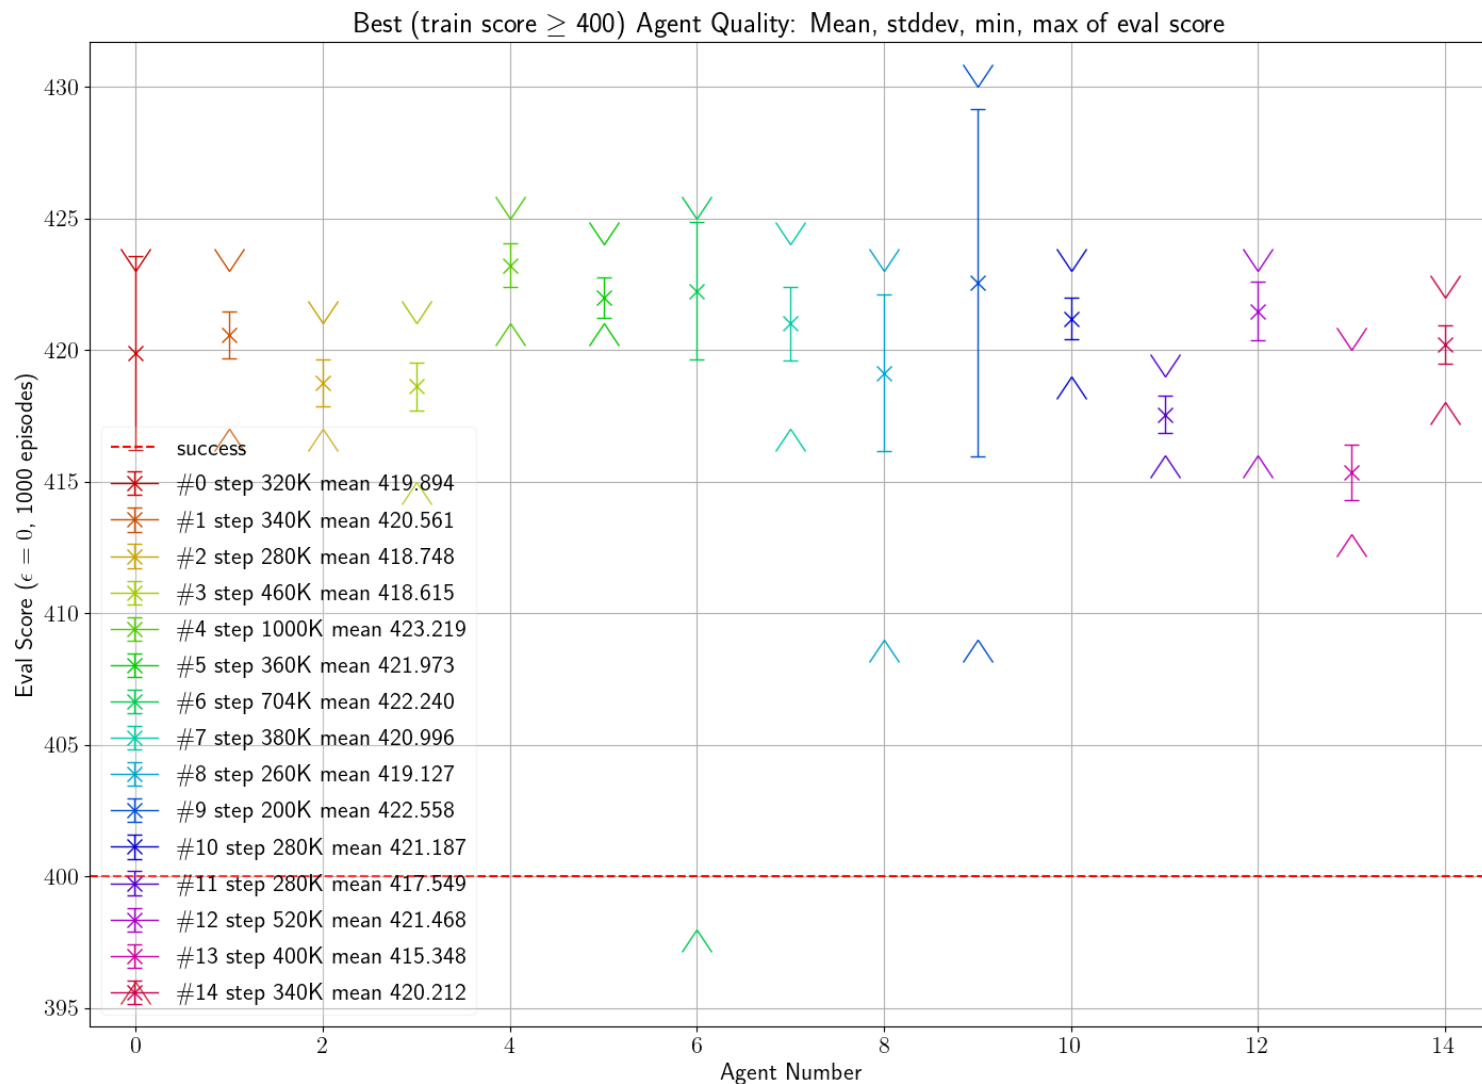


Figure 10: Quality test of best config with train metric

# Epsilon decay scheduler

Agents are already converging at high  $\epsilon$  values

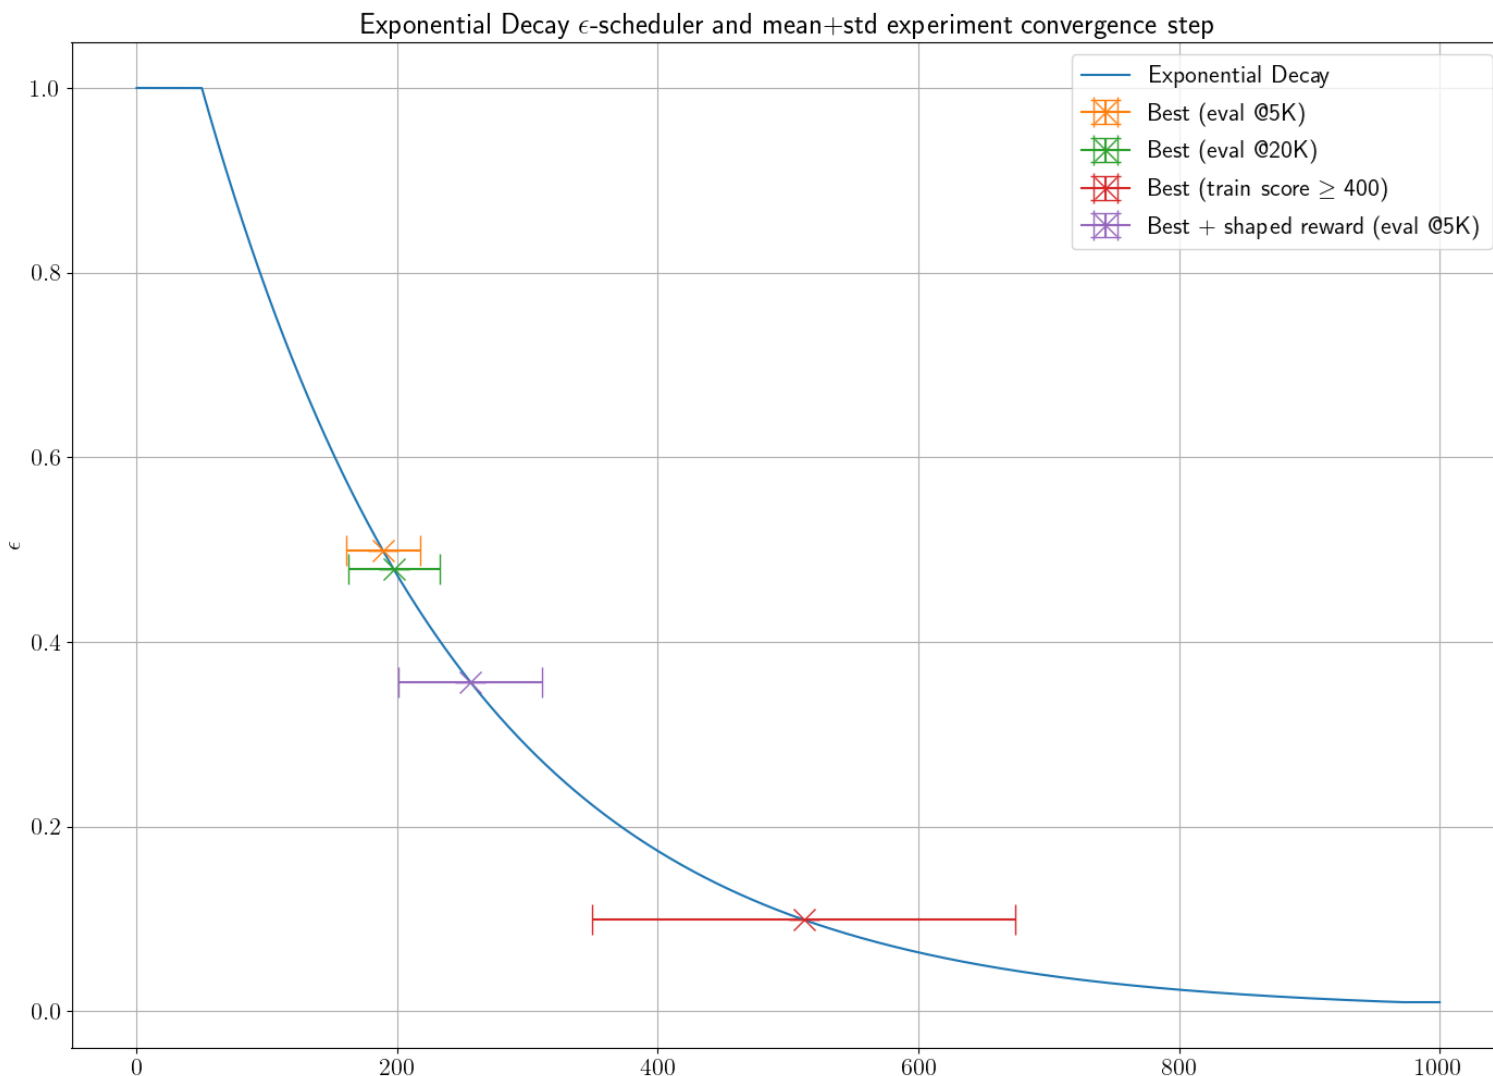


Figure 11: Convergence speeds and epsilon values of 4 selected experiments

# Runtime Analysis

Experiment	K steps	Episodes	Approx CPU Time	Approx GPU Time	Success metric	Converged / Total
Baseline	688.07±161.59	2851.13±714.40	799.62±187.79	636.84±149.56	Eval ≥ 400 (every 20K steps)	10/15
Best (eval @5K)	<b>189.02±28.41</b>	<b>879.73±80.56</b>	<b>585.58±88.03</b>	<b>284.59±42.78</b>	Eval ≥ 400 (every 5K steps)	15/15
Best (eval @20K)	197.41±35.29	893.93±110.19	606.49±108.42	297.52±53.19	Eval ≥ 400 (every 20K steps)	15/15
Best (train score)	512.08±162.60	1521.27±210.33	1584.34±503.07	758.81±240.94	Score of last 100 episodes ≥ 400	14/15
Best (shaped reward)	256.03±54.99	1027.00±163.02	720.96±154.84	326.83±70.19	Eval ≥ 400 (every 5K steps)	15/15

Table 2: Runtime Analysis of 5 selected experiments

- Runtime approximation process:
  - Run each experiment 5 times for 10000 steps
  - Calculate mean time per step
  - Multiply K Steps with mean time per step

# Configuration space for HPO

Underlined values: Baseline

**Bold** values: best configuration (lowest steps to convergence)

- Success metric: Eval every 20k steps, **eval every 5k steps**, avg train score last 100 ep
- Quantizer: how many actions (**2**, 3, 5, 16)
- Batch size (64, 128, **256**, 512)
- Discount: **0.99**
- Reward: **sparse**, shaped
- Optimizer/LR: **Adam/5e-4**, SGD/1e-2, RMSProp/1e-2, AdamW/5e-4
- Epsilon scheduler (first 5% steps  $\epsilon=1$ ): **Exponential Decay**, Cosine Decay
- DQN: (baseline, double Q, **dueling**, both)
- Tau (soft update): 1e-3, **2e-3**
- Replay Buffer Size: 100000, **500000**
- Nonlinearity: **ReLU**, SELU, LeakyReLU
- Regularization: Batchnorm, **None**
- Network Linear Layers: [64, 64], [**256**, **256**]
- State Space Transformation: None, **sincos transform on pole angle**

# All experiments (only solved runs)

## Experiments sorted by mean convergence step




Figure 12: Convergence speed of all experiments (only converged runs)

# Best agent over all experiments

- Our best single agent has a reward of  $426.5 \pm 0.84$ 
  - Tested over 1000 episodes
  - See *\_media/agent.mp4* for the video

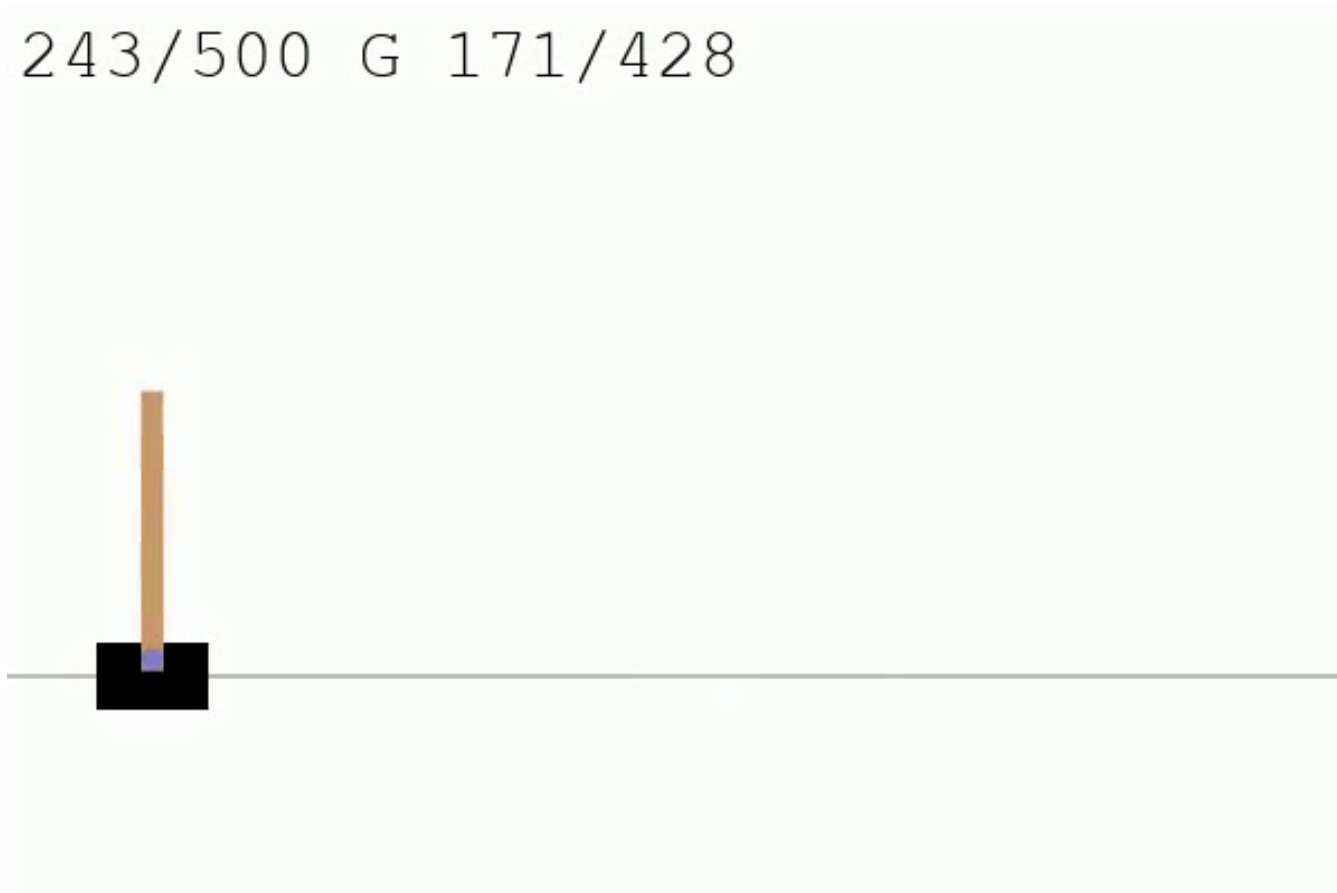


Figure 13: Best Agent in the Continuous Cartpole environment

# Qualitative analysis on other environments #1

- Our algorithm solves the Lunar Lander problem
  - 520K steps (~2400 episodes)
  - See `_media/lander.mp4` for the video




Figure 14: DQN Agent in the Lunar Lander environment

# Qualitative analysis on other environments #2

- Bipedal Walker is not solved after 1M steps
  - More punishment for applying stronger torque to the joints
  - Action space is 4-dimensional and continuous
  - Discretization is simply a bad strategy here
  - See *\_media/walker.mp4* for the video




Figure 15: DQN Agent in the Bipedal Walker environment



# Outlook, possible improvements

- Rainbow DQN<sup>[5]</sup> methods to improve the algorithm
- Improve epsilon decay schedule
  - Good for baseline, possibly too high  $\epsilon$  for better configs
  - Try Reward Based Epsilon Decay<sup>[6]</sup>
- Stabilize the train success metric experiments
  - Possibly related to the epsilon scheduler
  - Optimizer tuning / LR decay may be helpful
- Improve the eval success metric experiments
  - Use exponential moving average of weights
  - Increase time spent evaluating to get more accurate
- Automated HP search (e.g. Bayesian Optimization)
- Test more environments (Atari, classic control, ...)
- Compare results to other algorithms

# References

- [1] Mnih, V. et al (2015), Human Level Control Through Deep Reinforcement Learning  
<https://web.stanford.edu/class/psych209/Readings/MnihEtAlHassibis15NatureControlDeepRL.pdf>
- [2] Lillicrap, T. et al (2015), Continuous control with deep reinforcement learning  
<https://arxiv.org/abs/1509.02971>
- [3] van Hasselt, H. et al (2016), Deep Reinforcement Learning with Double Q-learning  
<https://arxiv.org/abs/1509.06461>
- [4] Wang, Z. et al (2015), Dueling Network Architectures for Deep Reinforcement Learning  
<https://arxiv.org/abs/1511.06581>
- [5] Hessel, M. et al (2017), Rainbow: Combining Improvements in Deep Reinforcement Learning  
<https://arxiv.org/abs/1710.02298>
- [6] Maroti, A., RBED: Reward Based Epsilon Decay  
<https://arxiv.org/abs/1910.13701>